

Depot


An infrastructure architect powered by AI,
from code to cloud.

TEAM

Ahsan Riaz
Muhammad Waleed
Muhammad Faizan Anwar

ADVISORS

Hira Anwar
Ayesha Hakim



Building apps is fast now. Shipping them is the bottleneck.

• Building is fast

~95%

of the code written by AI, at a quarter of YC's Winter 2025 start-ups

[TechCrunch, 2025](#) [1]

• Deploys break

~70%

of AI-written infrastructure failed to deploy on the first try, even from the best of six AI models

[Zhang et al., FSE 2026](#) [2]

• Costs run away

29%

of their cloud spend is wasted, say 753 cloud decision-makers

[Flexera State of the Cloud, 2026](#) [3]

• Few can ship it

17.8%

of developers work extensively with Terraform, against 43.3% on AWS

[Stack Overflow Developer Survey, 2025](#) [4]

Working code still leaves four decisions.

Decision 01



What it needs

Decision 02



Where it runs

Decision 03



What it costs

Decision 04



How to change it
safely

Depot answers all four, **and explains, records and approves every change.**

Each tool solves part of it.

Someone still decides where it runs.

Starts from the code

Picks where it runs

Asks before it acts

Managed platform



Vercel

— Plus platform config

× The platform only

— No approval gate

Bring-your-own-cloud



Flightcontrol

— Repo plus config file

— AWS only

— Deploy only

Proposed



Depot

✓ Ordinary repository

✓ Compares every option

✓ Signed, risk-tiered

Built for builders without an infrastructure engineer.



AI-assisted builders

A working product, but little experience of running one.



Solo full-stack developers

Outgrowing a managed platform: stay, or move to a cloud?



Small teams and agencies

Two to ten people, with no DevOps staff.

Three objectives.

What a user can do with Depot.

01

Deploy any supported repository

Point Depot at a repo. It reads the app, picks the target, prices it, provisions it.

02

Check health, metrics and cost

History, health, logs, spend and rollback in one place.

03

Move a live service, not just its code

Cutover planned, priced and rehearsed before the switch.

Research sits inside all three: comparing language models on extraction, and measured guardrails for the agent loop.

What Depot actually does.



Application Model

Every fact marked detected, inferred or user-supplied.



No stored keys

Bounded roles, federation, or one scoped token per target.



Architect

Picks among reviewed patterns, ranked by cost and effort.



Delivery and operations

Deploy on push, rollback, logs, cost, alerts, savings.



Safety gate

Validate, saved plan, policy, cost and role check.



Agent and incident loop

Changes by conversation; alarms end in a fix a person approves.

One session, from working code to live.

STEP 1

Connect

GitHub, GitLab or an
archive

STEP 2

Answer

Up to eight
questions

STEP 3

Review

Target, architecture,
monthly cost

STEP 4

Approve

Signed and role-
checked

STEP 5

Run

Deploy on push, roll
back, watch cost


What Depot covers.

IN SCOPE

- ✓ GitHub, GitLab or an archive
- ✓ Apps, APIs, static sites, workers
- ✓ AWS, GCP, Azure
- ✓ Short-lived credentials
- ✓ Approval before every change

Extended scope only if M8 is met early: guided mode, Oracle and Yandex Cloud, voice, preview environments.

Three clouds first, then everything else.



Team planning targets; dates will follow the official defence schedule.

The Architects

Three owners. No member reviews only their own work.



Ahsan Riaz

Application Model, agent mode,
incident loop, research

01

02

03




Muhammad Waleed

User-facing work, team workspaces,
managed platforms

01

02



Muhammad Faizan Anwar

Architect, cloud adapters, safety gate,
savings

01

03

Advisor **Hira Anwar** · Co-advisor **Ayesha Hakim**

Two goals we contribute to.

8 DECENT WORK AND
ECONOMIC GROWTH



SDG 8.3

Small-enterprise growth

Small teams run products on cloud accounts they own, at a known monthly cost, without hiring.

9 INDUSTRY, INNOVATION
AND INFRASTRUCTURE



SDG 9.5

Research and technological capability

An open benchmark, and views that explain why each architecture was chosen.

Sources.

1

[I. Mehta, "A quarter of startups in YC's current cohort have codebases that are almost entirely AI-generated," TechCrunch, Mar. 2025.](#)

2

[T. Zhang, S. Pan, Z. Zhang, Z. Xing and X. Sun, "Deployability-centric infrastructure-as-code generation: Fail, learn, refine, and succeed through LLM-empowered DevOps simulation," ACM FSE, 2026 \(arXiv:2506.05623\).](#)

3

[Flexera, 2026 State of the Cloud Report, 2026.](#)

4

[Stack Overflow, "Technology," 2025 Stack Overflow Developer Survey, 2025.](#)

5

[Vendor documentation of the nine products compared, accessed Sep. 30, 2026.](#)

6

[OWASP Gen AI Security Project, "LLM06:2025 Excessive agency," 2025.](#)

7

[United Nations, "Goal 8" and "Goal 9," Sustainable Development Goals, accessed Sep. 30, 2026.](#)

Click a source to open it. Full IEEE entries, with access dates, are in the proposal.

Thank you. Any questions?

• TEAM

Ahsan Riaz

Muhammad Waleed

Muhammad Faizan Anwar

ADVISORS

Hira Anwar

Ayesha Hakim

NUST SEECS · BS Computer Science